

Decision Trees  
oooooo

Regression Trees  
oooooooooooooooooooo

Pruning  
ooooooo

Trees in R  
oooooooooooo

# Regression Trees

Prof Wells

STA 295: Stat Learning

April 16th, 2024

## Outline

- Introduction to Decision Trees
- Discuss Theory and Algorithm for Decision Trees
- Describe the Pruning Algorithm as means of improving RMSE
- Implement Decision Trees in R

## Section 1

### Decision Trees

## Decision Trees

## Regression Trees

## Pruning

## Trees in R

My favorite animal is ...

My favorite animal is ...

Let's start with a game.

My favorite animal is . . .

Let's start with a game.

As a class, you may ask me up to **six** yes-or-no questions which I will truthfully answer, in order to learn more about my favorite animal.

After 6 questions, you may submit your guess for my favorite animal on a slip of paper.

- Don't look ahead on the slides!

My favorite animal is . . .

Let's start with a game.

As a class, you may ask me up to **six** yes-or-no questions which I will truthfully answer, in order to learn more about my favorite animal.

After 6 questions, you may submit your guess for my favorite animal on a slip of paper.

- Don't look ahead on the slides!

This slide intentionally blank

Don't look ahead at the next slide

My favorite animal is . . . the Pacific Tree Frog



## Guessing Algorithm

- The guessing algorithm determines a (partial) decision tree for classification

## Guessing Algorithm

- The guessing algorithm determines a (partial) decision tree for classification
  - Each Yes/No questions represent a branching point or node
  - The final guess represents the prediction
  - The 6 question limit restricts the depth of the tree
  - I am the test observation
  - What is the training data?

## Guessing Algorithm

- The guessing algorithm determines a (partial) decision tree for classification
  - Each Yes/No questions represent a branching point or node
  - The final guess represents the prediction
  - The 6 question limit restricts the depth of the tree
  - I am the test observation
  - What is the training data?
- What makes an effective question?

## Guessing Algorithm

- The guessing algorithm determines a (partial) decision tree for classification
  - Each Yes/No questions represent a branching point or node
  - The final guess represents the prediction
  - The 6 question limit restricts the depth of the tree
  - I am the test observation
  - What is the training data?
- What makes an effective question?
  - Separates data into roughly equal sizes
  - Data in each group are relatively similar
  - Later questions should be based on answers to earlier questions.
  - Early questions are general, later questions are specific.

Decision Trees  
ooooo

Regression Trees  
●ooooooooooooooo

Pruning  
ooooooo

Trees in R  
oooooooooooo

## Section 2

### Regression Trees

Decision Trees  
ooooo

Regression Trees  
○●oooooooooooooo

Pruning  
ooooooo

Trees in R  
oooooooooooo

## Regression Trees

Basic regression trees partition data into smaller groups that are more homogeneous with respect to predictors.

## Regression Trees

Basic regression trees partition data into smaller groups that are more homogeneous with respect to predictors.

- They then make predictions based on average value of response in each group

## Regression Trees

Basic regression trees partition data into smaller groups that are more homogeneous with respect to predictors.

- They then make predictions based on average value of response in each group

The most common technique is the Classification and Regression Tree (CART) method, which uses **Recursive Binary Splitting**

## Regression Trees

Basic regression trees partition data into smaller groups that are more homogeneous with respect to predictors.

- They then make predictions based on average value of response in each group

The most common technique is the Classification and Regression Tree (CART) method, which uses **Recursive Binary Splitting**

- ① The method begins with the entire data set  $S$  and searches every value of every predictor to cut  $S$  into two groups  $S_1$  and  $S_2$  that minimizes sum of squared error:

$$\text{SSE} = \sum_{i \in S_1} (y_i - \bar{y}_1)^2 + \sum_{i \in S_2} (y_i - \bar{y}_2)^2$$

## Regression Trees

Basic regression trees partition data into smaller groups that are more homogeneous with respect to predictors.

- They then make predictions based on average value of response in each group

The most common technique is the Classification and Regression Tree (CART) method, which uses **Recursive Binary Splitting**

- ① The method begins with the entire data set  $S$  and searches every value of every predictor to cut  $S$  into two groups  $S_1$  and  $S_2$  that minimizes sum of squared error:

$$\text{SSE} = \sum_{i \in S_1} (y_i - \bar{y}_1)^2 + \sum_{i \in S_2} (y_i - \bar{y}_2)^2$$

- ② The method then repeats step 1 for each of the two groups  $S_1$  and  $S_2$ .

## Regression Trees

Basic regression trees partition data into smaller groups that are more homogeneous with respect to predictors.

- They then make predictions based on average value of response in each group

The most common technique is the Classification and Regression Tree (CART) method, which uses **Recursive Binary Splitting**

- ① The method begins with the entire data set  $S$  and searches every value of every predictor to cut  $S$  into two groups  $S_1$  and  $S_2$  that minimizes sum of squared error:

$$\text{SSE} = \sum_{i \in S_1} (y_i - \bar{y}_1)^2 + \sum_{i \in S_2} (y_i - \bar{y}_2)^2$$

- ② The method then repeats step 1 for each of the two groups  $S_1$  and  $S_2$ .
- ③ The method continues splitting groups until each subdivision has few observations

## Regression Trees

Basic regression trees partition data into smaller groups that are more homogeneous with respect to predictors.

- They then make predictions based on average value of response in each group

The most common technique is the Classification and Regression Tree (CART) method, which uses **Recursive Binary Splitting**

- ① The method begins with the entire data set  $S$  and searches every value of every predictor to cut  $S$  into two groups  $S_1$  and  $S_2$  that minimizes sum of squared error:

$$\text{SSE} = \sum_{i \in S_1} (y_i - \bar{y}_1)^2 + \sum_{i \in S_2} (y_i - \bar{y}_2)^2$$

- ② The method then repeats step 1 for each of the two groups  $S_1$  and  $S_2$ .
  - ③ The method continues splitting groups until each subdivision has few observations.
- This is a *greedy* algorithm similar to forward selection, making the best choice at each stage. But it's not necessary that the algorithm creates a model with the best RSS
    - It's possible a suboptimal choice early could lead to extremely beneficial choice later, reducing the overall RSS

Decision Trees  
ooooo

Regression Trees  
oo●oooooooooooo

Pruning  
ooooooo

Trees in R  
oooooooooooo

# Trees

Portland, OR is known for its coffee, its politics, its 117 degree summers, but also its ...

Decision Trees  
oooooo

Regression Trees  
oo●ooooooooooooo

Pruning  
ooooooo

Trees in R  
oooooooooooo

# Trees

Portland, OR is known for its coffee, its politics, its 117 degree summers, but also its ...

Trees!



Decision Trees  
ooooo

Regression Trees  
ooo●oooooooooooo

Pruning  
ooooooo

Trees in R  
oooooooooooo

## pdxTrees

We'll study regression trees using a subset of the pdxTrees data set from the pdxTrees GitHub repo (Maintained by K. McConville, I. Caldwell, and N. Horton)

## pdxTrees

We'll study regression trees using a subset of the pdxTrees data set from the pdxTrees GitHub repo (Maintained by K. McConville, I. Caldwell, and N. Horton)

- The data was collected by the Portland Parks and Rec's *Urban Forestry Tree Inventory Project*.
- The Tree Inventory Project has gathered data on Portland trees since 2010, collecting this data in the summer months with a team of over 1,300 volunteers and city employees.

## pdxTrees

We'll study regression trees using a subset of the pdxTrees data set from the pdxTrees GitHub repo (Maintained by K. McConville, I. Caldwell, and N. Horton)

- The data was collected by the Portland Parks and Rec's *Urban Forestry Tree Inventory Project*.
- The Tree Inventory Project has gathered data on Portland trees since 2010, collecting this data in the summer months with a team of over 1,300 volunteers and city employees.
- The pdxTrees dataset is too large to install alongside the package. Instead, the package provides helper loading functions:
  - `get_pdxTrees_parks()` pulls data on 25,534 trees from 174 Portland parks
  - `get_pdxTrees_streets()` pulls data on 218,602 trees along Portland streets

## pdxTrees Data

- To keep things manageable, we'll limit our study to trees in just a few parks:

```
library(pdxTrees)
my_pdxTrees <- get_pdxTrees_parks(park = c("Kenilworth Park", "Westmoreland Park",
                                             "Woodstock Park", "Berkeley Park", "Powell Park"))
```

## pdxTrees Data

- To keep things manageable, we'll limit our study to trees in just a few parks:

```
library(pdxTrees)
my_pdxTrees <- get_pdxTrees_parks(park = c("Kenilworth Park", "Westmoreland Park",
                                             "Woodstock Park", "Berkeley Park", "Powell Park"))
```

- How many observations?

```
dim(my_pdxTrees)
## [1] 1039   35
```

## pdxTrees Data

- To keep things manageable, we'll limit our study to trees in just a few parks:

```
library(pdxTrees)
my_pdxTrees <- get_pdxTrees_parks(park = c("Kenilworth Park", "Westmoreland Park",
                                             "Woodstock Park", "Berkeley Park", "Powell Park"))
```

- How many observations?

```
dim(my_pdxTrees)
```

```
## [1] 1039    35
```

- What variables are present?

```
names(my_pdxTrees)
```

## pdxTrees Data

- To keep things manageable, we'll limit our study to trees in just a few parks:

```
library(pdxTrees)
my_pdxTrees <- get_pdxTrees_parks(park = c("Kenilworth Park", "Westmoreland Park",
                                             "Woodstock Park", "Berkeley Park", "Powell Park"))
```

- How many observations?

```
dim(my_pdxTrees)
```

```
## [1] 1039    35
```

- What variables are present?

```
names(my_pdxTrees)
```

```
##  [1] "Longitude"          "Latitude"
##  [3] "UserID"              "Genus"
##  [5] "Family"              "DBH"
##  [7] "Inventory_Date"      "Species"
##  [9] "Common_Name"         "Condition"
## [11] "Tree_Height"         "Crown_Width_NS"
## [13] "Crown_Width_EW"      "Crown_Base_Height"
## [15] "Collected_By"        "Park"
## [17] "Scientific_Name"      "Functional_Type"
## [19] "Mature_Size"         "Native"
## [21] "Edible"               "Nuisance"
## [23] "Structural_Value"     "Carbon_Storage_lb"
## [25] "Carbon_Storage_value" "Carbon_Sequestration_lb"
## [27] "Carbon_Sequestration_value" "Stormwater_ft"
## [29] "Stormwater_value"      "Pollution_Removal_value"
## [31] "Pollution_Removal_oz"   "Total_Annual_Services"
## [33] "Origin"                "Species_Factoid"
## [35] "Crown_Width"
```

Decision Trees  
ooooo

Regression Trees  
oooooo●oooooooooooo

Pruning  
ooooooo


Trees in R  
oooooooooooo

## Carbon Sequestration

- Forestry carbon sequestration is the process by which trees capture and store carbon dioxide from the atmosphere


# Carbon Sequestration

- Forestry carbon sequestration is the process by which trees capture and store carbon dioxide from the atmosphere



# Carbon Sequestration


- Forestry carbon sequestration is the process by which trees capture and store carbon dioxide from the atmosphere



- Annual carbon sequestration of tree depends on several factors:

# Carbon Sequestration


- Forestry carbon sequestration is the process by which trees capture and store carbon dioxide from the atmosphere



- Annual carbon sequestration of tree depends on several factors:
  - Species, Size, Age, Location, Weather, etc.

# Carbon Sequestration

- Forestry carbon sequestration is the process by which trees capture and store carbon dioxide from the atmosphere



- Annual carbon sequestration of tree depends on several factors:
  - Species, Size, Age, Location, Weather, etc.
- Who might be interested in estimating the carbon sequestration of a tree?
  - Why?

## Predicting Carbon Sequestration

- Can we predict carbon sequestration based on other tree features?
  - In pdxTrees, annual carbon sequestration is encoded as Carbon\_Sequestration\_lb.

## Predicting Carbon Sequestration

- Can we predict carbon sequestration based on other tree features?
  - In pdxTrees, annual carbon sequestration is encoded as Carbon\_Sequestration\_lb.
- Let's focus on size-related variables:
  - Tree\_height, height from ground to the live top of tree (in ft)
  - Crown\_Width, the average of North-South and East-West canopy width (in ft)

## Predicting Carbon Sequestration

- Can we predict carbon sequestration based on other tree features?
  - In pdxTrees, annual carbon sequestration is encoded as Carbon\_Sequestration\_lb.
- Let's focus on size-related variables:
  - Tree\_height, height from ground to the live top of tree (in ft)
  - Crown\_Width, the average of North-South and East-West canopy width (in ft)
- We are looking at a model of the form

Carbon\_Sequestration ~ Tree\_Height + Crown\_Width

## Predicting Carbon Sequestration

- Can we predict carbon sequestration based on other tree features?
  - In pdxTrees, annual carbon sequestration is encoded as Carbon\_Sequestration\_lb.
- Let's focus on size-related variables:
  - Tree\_height, height from ground to the live top of tree (in ft)
  - Crown\_Width, the average of North-South and East-West canopy width (in ft)
- We are looking at a model of the form

Carbon\_Sequestration ~ Tree\_Height + Crown\_Width

- What are correlations among these 3 variables?

## Predicting Carbon Sequestration

- Can we predict carbon sequestration based on other tree features?
  - In pdxTrees, annual carbon sequestration is encoded as Carbon\_Sequestration\_lb.
- Let's focus on size-related variables:
  - Tree\_height, height from ground to the live top of tree (in ft)
  - Crown\_Width, the average of North-South and East-West canopy width (in ft)
- We are looking at a model of the form





Carbon\_Sequestration ~ Tree\_Height + Crown\_Width

- What are correlations among these 3 variables?

```
##                                     Carbon_Sequestration_lb Crown_Width Tree_Height
## Carbon_Sequestration_lb               1.0000000   0.6126951   0.4362020
## Crown_Width                         0.6126951   1.0000000   0.5980118
## Tree_Height                        0.4362020   0.5980118   1.0000000
```





## Predicting Carbon Sequestration

- Can we predict carbon sequestration based on other tree features?



## Predicting Carbon Sequestration

- Can we predict carbon sequestration based on other tree features?



- Observations?

Decision Trees  
oooooo

Regression Trees  
oooooooo●ooooo

Pruning  
ooooooo

Trees in R  
oooooooooooo

## An Old Friend

This seems like a good time to implement linear regression:

# An Old Friend

This seems like a good time to implement linear regression:

```
tree_lm<-lm(Carbon_Sequestration_lb ~Crown_Width + Tree_Height, data=my_pdxTrees)
summary(tree_lm)
```

```
##
## Call:
## lm(formula = Carbon_Sequestration_lb ~ Crown_Width + Tree_Height,
##      data = my_pdxTrees)
##
## Residuals:
##    Min      1Q  Median      3Q     Max
## -87.395 -13.283   4.912  10.982 121.950
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) -3.08819   2.03721 -1.516 0.129853
## Crown_Width  0.88769   0.04947 17.944 < 2e-16 ***
## Tree_Height  0.10140   0.02848  3.560 0.000388 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 26.46 on 1031 degrees of freedom
## (5 observations deleted due to missingness)
## Multiple R-squared:  0.383, Adjusted R-squared:  0.3818
## F-statistic:  320 on 2 and 1031 DF, p-value: < 2.2e-16
```


Decision Trees  
oooooo

Regression Trees  
oooooooooooo●oooo


Pruning  
ooooooo

Trees in R  
oooooooooooo

## Diagnostic Plots



## Diagnostic Plots



- Concerns?

Decision Trees  
oooooo

Regression Trees  
oooooooooooo●ooo

Pruning  
ooooooo


Trees in R  
oooooooooooo

## Regression Tree

- Instead, let's build a regression tree:

## Regression Tree

- Instead, let's build a regression tree:



- Leaves at the bottom of the tree provide predictions


Decision Trees  
oooooo

Regression Trees  
oooooooooooooooooooo

Pruning  
ooooooo

Trees in R  
oooooooooooo

## Another Visualization



## Interpretation

- Crown\_Width is the most important predictor of Carbon\_Sequestration\_lb
- After accounting for width, Tree\_Height has some impact on Carbon\_Sequestration\_lb
- Very narrow and very wide trees tend to have low Carbon\_Sequestration\_lb
- Trees of moderate width and height have largest Carbon\_Sequestration\_lb

## Tree Accuracy

- Let's create a test set consisting of two other parks:

```
my_pdxTrees_test <- get_pdxTrees_parks(park = c("Mt Scott Park", "Glenwood Park"))
```

## Tree Accuracy

- Let's create a test set consisting of two other parks:

```
my_pdxTrees_test <- get_pdxTrees_parks(park = c("Mt Scott Park", "Glenwood Park"))
```

- We'll measure the accuracy of our model using root Mean Square Error:

$$\text{rMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

```
## Tree_rMSE
## 1 15.37258
```

## Tree Accuracy

- Let's create a test set consisting of two other parks:

```
my_pdxTrees_test <- get_pdxTrees_parks(park = c("Mt Scott Park", "Glenwood Park"))
```

- We'll measure the accuracy of our model using root Mean Square Error:

$$\text{rMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

```
## Tree_rMSE
## 1 15.37258
```

- And compared to the linear model:

```
## lm_rMSE
## 1 16.87355
```

## Tree Accuracy

- Let's create a test set consisting of two other parks:

```
my_pdxTrees_test <- get_pdxTrees_parks(park = c("Mt Scott Park", "Glenwood Park"))
```

- We'll measure the accuracy of our model using root Mean Square Error:

$$\text{rMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

```
## Tree_rMSE  
## 1 15.37258
```

- And compared to the linear model:

```
## lm_rMSE  
## 1 16.87355
```

- Why did the tree model outperform the linear model?

## Tree Accuracy

- Let's create a test set consisting of two other parks:

```
my_pdxTrees_test <- get_pdxTrees_parks(park = c("Mt Scott Park", "Glenwood Park"))
```

- We'll measure the accuracy of our model using root Mean Square Error:

$$\text{rMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

```
## Tree_rMSE  
## 1 15.37258
```

- And compared to the linear model:

```
## lm_rMSE  
## 1 16.87355
```

- Why did the tree model outperform the linear model?

- Nevertheless, what are some downsides to the tree model?

Decision Trees  
ooooo

Regression Trees  
oooooooooooooooooooo

Pruning  
●oooooo

Trees in R  
oooooooooooo

## Section 3

### Pruning

Decision Trees  
ooooo

Regression Trees  
oooooooooooooooooooo

Pruning  
○●ooooo

Trees in R  
oooooooooooo

## Model Accuracy and Tree Algorithm

Recall the CART algorithm:

# Model Accuracy and Tree Algorithm

Recall the CART algorithm:

- ① Cut the data set  $S$  into two groups  $S_1$  and  $S_2$  that minimizes sum of squared error:

$$\text{SSE} = \sum_{i \in S_1} (y_i - \bar{y}_1)^2 + \sum_{i \in S_2} (y_i - \bar{y}_2)^2$$

- ② Repeats step 1 for each of the two groups  $S_1$  and  $S_2$ .
- ③ Continue splitting groups until each subdivision has few observations

## Model Accuracy and Tree Algorithm

Recall the CART algorithm:

- ① Cut the data set  $S$  into two groups  $S_1$  and  $S_2$  that minimizes sum of squared error:

$$\text{SSE} = \sum_{i \in S_1} (y_i - \bar{y}_1)^2 + \sum_{i \in S_2} (y_i - \bar{y}_2)^2$$

- ② Repeats step 1 for each of the two groups  $S_1$  and  $S_2$ .
  - ③ Continue splitting groups until each subdivision has few observations
- This model will tend to produce accurate estimates on the *training set* (why?)

# Model Accuracy and Tree Algorithm

Recall the CART algorithm:

- ① Cut the data set  $S$  into two groups  $S_1$  and  $S_2$  that minimizes sum of squared error:

$$\text{SSE} = \sum_{i \in S_1} (y_i - \bar{y}_1)^2 + \sum_{i \in S_2} (y_i - \bar{y}_2)^2$$

- ② Repeats step 1 for each of the two groups  $S_1$  and  $S_2$ .
  - ③ Continue splitting groups until each subdivision has few observations
- This model will tend to produce accurate estimates on the *training set* (why?)
  - However, it is extremely liable to overfit, and produce inaccurate estimates on *test set*

## Model Accuracy and Tree Algorithm

Recall the CART algorithm:

- ① Cut the data set  $S$  into two groups  $S_1$  and  $S_2$  that minimizes sum of squared error:

$$\text{SSE} = \sum_{i \in S_1} (y_i - \bar{y}_1)^2 + \sum_{i \in S_2} (y_i - \bar{y}_2)^2$$

- ② Repeats step 1 for each of the two groups  $S_1$  and  $S_2$ .
  - ③ Continue splitting groups until each subdivision has few observations
- This model will tend to produce accurate estimates on the *training set* (why?)
  - However, it is extremely liable to overfit, and produce inaccurate estimates on *test set*
  - To improve test RMSE, we need to reduce variance in model estimates (at the potential cost of some increased bias)

## Model Accuracy and Tree Algorithm

Recall the CART algorithm:

- ① Cut the data set  $S$  into two groups  $S_1$  and  $S_2$  that minimizes sum of squared error:

$$\text{SSE} = \sum_{i \in S_1} (y_i - \bar{y}_1)^2 + \sum_{i \in S_2} (y_i - \bar{y}_2)^2$$

- ② Repeats step 1 for each of the two groups  $S_1$  and  $S_2$ .
  - ③ Continue splitting groups until each subdivision has few observations
- This model will tend to produce accurate estimates on the *training set* (why?)
  - However, it is extremely liable to overfit, and produce inaccurate estimates on *test set*
  - To improve test RMSE, we need to reduce variance in model estimates (at the potential cost of some increased bias)
  - To accomplish this, we can force trees to be less thorough in creating homogeneous groups (i.e. have fewer nodes)

## Model Accuracy and Tree Algorithm

Recall the CART algorithm:

- ① Cut the data set  $S$  into two groups  $S_1$  and  $S_2$  that minimizes sum of squared error:

$$\text{SSE} = \sum_{i \in S_1} (y_i - \bar{y}_1)^2 + \sum_{i \in S_2} (y_i - \bar{y}_2)^2$$

- ② Repeats step 1 for each of the two groups  $S_1$  and  $S_2$ .
  - ③ Continue splitting groups until each subdivision has few observations
- This model will tend to produce accurate estimates on the *training set* (why?)
  - However, it is extremely liable to overfit, and produce inaccurate estimates on *test set*
  - To improve test RMSE, we need to reduce variance in model estimates (at the potential cost of some increased bias)
  - To accomplish this, we can force trees to be less thorough in creating homogeneous groups (i.e. have fewer nodes)
    - Option 1: Grow “younger” trees that are shorter
    - Option 2: Grow “mature” trees that are longer, and prune them back

Decision Trees  
oooooo

Regression Trees  
oooooooooooooooooooo

Pruning  
oo●oooo

Trees in R  
oooooooooooo

## Subtrees

Recall that when we grow trees, we use the greedy recursive binary splitting algorithm

## Subtrees

Recall that when we grow trees, we use the greedy recursive binary splitting algorithm

- This means any splits that are made early will persist throughout the rest of the model

## Subtrees

Recall that when we grow trees, we use the greedy recursive binary splitting algorithm

- This means any splits that are made early will persist throughout the rest of the model
- For this reason, growing a “young” tree (Method 1) tends to be ineffective, since the model may be burdened with a poor early choice

## Subtrees

Recall that when we grow trees, we use the greedy recursive binary splitting algorithm

- This means any splits that are made early will persist throughout the rest of the model
- For this reason, growing a “young” tree (Method 1) tends to be ineffective, since the model may be burdened with a poor early choice
- The better strategy is to grow a full tree, and then remove some of the less helpful terminal leaves (Method 2)

## Subtrees

Recall that when we grow trees, we use the greedy recursive binary splitting algorithm

- This means any splits that are made early will persist throughout the rest of the model
- For this reason, growing a “young” tree (Method 1) tends to be ineffective, since the model may be burdened with a poor early choice
- The better strategy is to grow a full tree, and then remove some of the less helpful terminal leaves (Method 2)

A **subtree** is a regression tree obtained by removing some of the branches and nodes from the full regression tree.

## Subtrees

Recall that when we grow trees, we use the greedy recursive binary splitting algorithm

- This means any splits that are made early will persist throughout the rest of the model
- For this reason, growing a “young” tree (Method 1) tends to be ineffective, since the model may be burdened with a poor early choice
- The better strategy is to grow a full tree, and then remove some of the less helpful terminal leaves (Method 2)

A **subtree** is a regression tree obtained by removing some of the branches and nodes from the full regression tree.

- A subtree will tend to have higher *training* RSS than the full tree. But can often have lower *test* RSS than the full tree.

Decision Trees  
ooooo

Regression Trees  
oooooooooooooooooooo

Pruning  
ooooooo

Trees in R  
oooooooooooo

## Pruning Algorithm

- Once a tree is fully grown, we *prune* it using *cost-complexity tuning*

## Pruning Algorithm

- Once a tree is fully grown, we *prune* it using *cost-complexity tuning*
  - The goal is to find a tree of optimal size with the smallest error rate. We consider a sequence of trees indexed by a tuning parameter  $\alpha$ .

## Pruning Algorithm

- Once a tree is fully grown, we *prune* it using *cost-complexity tuning*
  - The goal is to find a tree of optimal size with the smallest error rate. We consider a sequence of trees indexed by a tuning parameter  $\alpha$ .
- For each value of  $\alpha$ , there exists a unique subtree  $T$  of the full tree  $T_0$  that minimizes

$$\text{RSS} + \alpha|T|$$

where  $|T|$  is the number of terminal nodes of the tree  $T$ .

## Pruning Algorithm

- Once a tree is fully grown, we *prune* it using *cost-complexity tuning*
  - The goal is to find a tree of optimal size with the smallest error rate. We consider a sequence of trees indexed by a tuning parameter  $\alpha$ .
- For each value of  $\alpha$ , there exists a unique subtree  $T$  of the full tree  $T_0$  that minimizes

$$\text{RSS} + \alpha|T|$$

where  $|T|$  is the number of terminal nodes of the tree  $T$ .

- That is,  $\alpha$  penalizes a tree based on its number of terminal nodes.

## Pruning Algorithm

- Once a tree is fully grown, we *prune* it using *cost-complexity tuning*
  - The goal is to find a tree of optimal size with the smallest error rate. We consider a sequence of trees indexed by a tuning parameter  $\alpha$ .
- For each value of  $\alpha$ , there exists a unique subtree  $T$  of the full tree  $T_0$  that minimizes

$$\text{RSS} + \alpha|T|$$

where  $|T|$  is the number of terminal nodes of the tree  $T$ .

- That is,  $\alpha$  penalizes a tree based on its number of terminal nodes.
- This is analogous to the penalty parameter in Penalized Regression (LASSO / Ridge Regression) as well as in Cp/AIC/BIC for Best Subset

## Pruning Algorithm

- Once a tree is fully grown, we *prune* it using *cost-complexity tuning*
  - The goal is to find a tree of optimal size with the smallest error rate. We consider a sequence of trees indexed by a tuning parameter  $\alpha$ .
- For each value of  $\alpha$ , there exists a unique subtree  $T$  of the full tree  $T_0$  that minimizes

$$\text{RSS} + \alpha|T|$$

where  $|T|$  is the number of terminal nodes of the tree  $T$ .

- That is,  $\alpha$  penalizes a tree based on its number of terminal nodes.
- This is analogous to the penalty parameter in Penalized Regression (LASSO / Ridge Regression) as well as in Cp/AIC/BIC for Best Subset
- As  $\alpha$  increases from 0 (i.e. the full tree), branches get pruned in a predictable way, making for relatively quick computation.

## Pruning Algorithm

- Once a tree is fully grown, we *prune* it using *cost-complexity tuning*
  - The goal is to find a tree of optimal size with the smallest error rate. We consider a sequence of trees indexed by a tuning parameter  $\alpha$ .
- For each value of  $\alpha$ , there exists a unique subtree  $T$  of the full tree  $T_0$  that minimizes

$$\text{RSS} + \alpha|T|$$

where  $|T|$  is the number of terminal nodes of the tree  $T$ .

- That is,  $\alpha$  penalizes a tree based on its number of terminal nodes.
- This is analogous to the penalty parameter in Penalized Regression (LASSO / Ridge Regression) as well as in Cp/AIC/BIC for Best Subset
- As  $\alpha$  increases from 0 (i.e. the full tree), branches get pruned in a predictable way, making for relatively quick computation.
- We can find the optimal value of  $\alpha$  by further splitting training data into a training and test set (or using k-fold cross-validation). We can choose the **best** subtree by:

## Pruning Algorithm

- Once a tree is fully grown, we *prune* it using *cost-complexity tuning*
  - The goal is to find a tree of optimal size with the smallest error rate. We consider a sequence of trees indexed by a tuning parameter  $\alpha$ .
- For each value of  $\alpha$ , there exists a unique subtree  $T$  of the full tree  $T_0$  that minimizes

$$\text{RSS} + \alpha|T|$$

where  $|T|$  is the number of terminal nodes of the tree  $T$ .

- That is,  $\alpha$  penalizes a tree based on its number of terminal nodes.
- This is analogous to the penalty parameter in Penalized Regression (LASSO / Ridge Regression) as well as in Cp/AIC/BIC for Best Subset
- As  $\alpha$  increases from 0 (i.e. the full tree), branches get pruned in a predictable way, making for relatively quick computation.
- We can find the optimal value of  $\alpha$  by further splitting training data into a training and test set (or using k-fold cross-validation). We can choose the **best** subtree by:
  - Choosing the tree with smallest rMSE.

## Pruning Algorithm

- Once a tree is fully grown, we *prune* it using *cost-complexity tuning*
  - The goal is to find a tree of optimal size with the smallest error rate. We consider a sequence of trees indexed by a tuning parameter  $\alpha$ .
- For each value of  $\alpha$ , there exists a unique subtree  $T$  of the full tree  $T_0$  that minimizes


$$\text{RSS} + \alpha|T|$$

where  $|T|$  is the number of terminal nodes of the tree  $T$ .

- That is,  $\alpha$  penalizes a tree based on its number of terminal nodes.
  - This is analogous to the penalty parameter in Penalized Regression (LASSO / Ridge Regression) as well as in Cp/AIC/BIC for Best Subset
- As  $\alpha$  increases from 0 (i.e. the full tree), branches get pruned in a predictable way, making for relatively quick computation.
  - We can find the optimal value of  $\alpha$  by further splitting training data into a training and test set (or using k-fold cross-validation). We can choose the **best** subtree by:
    - Choosing the tree with smallest rMSE.
    - Choosing the *smallest* tree with rMSE within 1 standard deviation of lowest rMSE

## Pruning Example


- How does rMSE vary as tree size changes?



- Horizontal axis gives values of complexity parameter (cp)
- Upper scale indicates number of terminal nodes for given tree
- Vertical axis gives the cross-validated relative root mean squared error
- Dotted horizontal line has height equal to 1 standard error above smallest rMSE

## Pruning Example


- How does rMSE vary as tree size changes?



- What are the test MSEs for the full tree and the subtrees with 4 and 8 leaves?

```
## # A tibble: 4 x 4
##   model      .metric .estimator .estimate
##   <chr>      <chr>   <chr>        <dbl>
## 1 pruned     rmse    standard     14.3 
## 2 full       rmse    standard     15.4 
## 3 linear     rmse    standard     16.9 
## 4 very pruned rmse    standard     16.9 
```

# Comparison



Decision Trees  
ooooo

Regression Trees  
oooooooooooooooooooo

Pruning  
ooooooo

Trees in R  
●oooooooooooo

## Section 4

### Trees in R

Decision Trees  
ooooo

Regression Trees  
oooooooooooooooooooo

Pruning  
ooooooo

Trees in R  
o●oooooooooooo

## Creating Tree Models in R

There are two common packages for creating regression trees in R: `tree` and `rpart`.

## Creating Tree Models in R

There are two common packages for creating regression trees in R: `tree` and `rpart`.

- The `tree` package is one of the oldest packages on CRAN. It is a (tiny) bit easier to use. But allows far less customization. (Traditional)

## Creating Tree Models in R

There are two common packages for creating regression trees in R: `tree` and `rpart`.

- The `tree` package is one of the oldest packages on CRAN. It is a (tiny) bit easier to use. But allows far less customization. (Traditional)
- The `rpart` package is newer, computationally faster, and has more options. It also can be combined with other packages for **much** nicer plots. (Recommended)

## Trees using 'rpart'

- To fit a tree using variables Tree\_Height and Crown\_Width:

```
set.seed(1)
library(rpart)
tree_model1 <- rpart(Carbon_Sequestration_lb ~
                      Tree_Height + Crown_Width,
                      data = my_pdxTrees)
```

## Trees using 'rpart'

- To fit a tree using variables Tree\_Height and Crown\_Width:

```
set.seed(1)
library(rpart)
tree_model1 <- rpart(Carbon_Sequestration_lb ~
                      Tree_Height + Crown_Width,
                      data = my_pdxTrees)
```

- We can change several features of the tree by adding a control argument:

```
set.seed(1)
tree_model2 <- rpart(Carbon_Sequestration_lb ~
                      Tree_Height + Crown_Width,
                      control = rpart.control(
                        minsplit = 20, xval = 10, maxdepth = 10, cp = 0.005),
                      data = my_pdxTrees)
```

## Trees using 'rpart'

- To fit a tree using variables Tree\_Height and Crown\_Width:

```
set.seed(1)
library(rpart)
tree_model1 <- rpart(Carbon_Sequestration_lb ~
                      Tree_Height + Crown_Width,
                      data = my_pdxTrees)
```

- We can change several features of the tree by adding a control argument:

```
set.seed(1)
tree_model2 <- rpart(Carbon_Sequestration_lb ~
                      Tree_Height + Crown_Width,
                      control = rpart.control(
                        minsplit = 20, xval = 10, maxdepth = 10, cp = 0.005),
                      data = my_pdxTrees)
```

- `minsplit` is the minimum number of observations in a node
- `xval` is the number of cross-validation folds used
- `maxdepth` is the maximum depth of any node in the final tree
- `cp` is the minimum reduction in RSS needed in order to attempt a split

Decision Trees  
ooooo

Regression Trees  
oooooooooooooooooooo

Pruning  
ooooooo

Trees in R  
ooo●oooooooo


## Plots using plot

- There are several options for visualizing trees with varying ease-of-use and aesthetics.
  - The base R `plot` function quickly generates plots, but...

## Plots using plot

- There are several options for visualizing trees with varying ease-of-use and aesthetics.
  - The base R plot function quickly generates plots, but...


```
plot(tree_model2)
text(tree_model2, pretty = 0, cex = .5)
```



## Plots using rpart.plot

- An alternative to plot is the rpart.plot function from the package of the same name:

```
library(rpart.plot)
rpart.plot(tree_model2)
```



- Some further customization available (see ?rpart.plot)

Decision Trees  
ooooo

Regression Trees  
oooooooooooooooooooo

Pruning  
ooooooo

Trees in R  
oooooo●oooo

## Trees in R via rpart cont'd

- The `rpart` function automatically performs  $k$ -fold CV when choosing among potential splits.

## Trees in R via rpart cont'd

- The `rpart` function automatically performs  $k$ -fold CV when choosing among potential splits.
- To access results, append `$cptable` to the `rpart` model object:

```
tree_model2$cptable
```

```
##          CP nsplit rel error      xerror       xstd
## 1  0.304594426      0 1.0000000 1.0015819 0.06083997
## 2  0.127260632      1 0.6954056 0.7296616 0.04066145
## 3  0.025587347      2 0.5681449 0.6138089 0.03803769
## 4  0.015177861      3 0.5425576 0.5753989 0.03914593
## 5  0.014222123      5 0.5122019 0.5705610 0.03960569
## 6  0.010849075      6 0.4979797 0.5548873 0.03897808
## 7  0.009024312      7 0.4871307 0.5342982 0.03851474
## 8  0.006608748      9 0.4690820 0.5269734 0.04016718
## 9  0.006064592     10 0.4624733 0.5365762 0.04102266
## 10 0.005000000     11 0.4564087 0.5360229 0.04159291
```

## Trees in R via rpart cont'd

- The `rpart` function automatically performs  $k$ -fold CV when choosing among potential splits.
- To access results, append `$cptable` to the `rpart` model object:

```
tree_model2$cptable
```

```
##          CP nsplit rel error      xerror       xstd
## 1  0.304594426      0 1.0000000 1.0015819 0.06083997
## 2  0.127260632      1 0.6954056 0.7296616 0.04066145
## 3  0.025587347      2 0.5681449 0.6138089 0.03803769
## 4  0.015177861      3 0.5425576 0.5753989 0.03914593
## 5  0.014222123      5 0.5122019 0.5705610 0.03960569
## 6  0.010849075      6 0.4979797 0.5548873 0.03897808
## 7  0.009024312      7 0.4871307 0.5342982 0.03851474
## 8  0.006608748      9 0.4690820 0.5269734 0.04016718
## 9  0.006064592     10 0.4624733 0.5365762 0.04102266
## 10 0.005000000     11 0.4564087 0.5360229 0.04159291
```

- CP is the value of the complexity parameter
- nsplit is number of splits
- rel error is  $(1 - R^2)$ , using  $R^2 = 1 - \frac{\text{RSS}}{\text{TSS}}$
- xerror is cross-validated estimate of relative error
- xstd is the standard deviation in xerror based on CV

## Analyze Results

- The printcp function displays key model information

```
printcp(tree_model2)
```

```
##  
## Regression tree:  
## rpart(formula = Carbon_Sequestration_lb ~ Tree_Height + Crown_Width,  
##       data = my_pdxTrees, control = rpart.control(minsplit = 20,  
##             xval = 10, maxdepth = 10, cp = 0.005))  
##  
## Variables actually used in tree construction:  
## [1] Crown_Width Tree_Height  
##  
## Root node error: 1175664/1037 = 1133.7  
##  
## n=1037 (2 observations deleted due to missingness)  
##  
##          CP nsplit rel error  xerror      xstd  
## 1  0.3045944     0   1.00000 1.00158  0.060840  
## 2  0.1272606     1   0.69541 0.72966  0.040661  
## 3  0.0255873     2   0.56814 0.61381  0.038038  
## 4  0.0151779     3   0.54256 0.57540  0.039146  
## 5  0.0142221     5   0.51220 0.57056  0.039606  
## 6  0.0108491     6   0.49798 0.55489  0.038978  
## 7  0.0090243     7   0.48713 0.53430  0.038515  
## 8  0.0066087     9   0.46908 0.52697  0.040167  
## 9  0.0060646    10   0.46247 0.53658  0.041023  
## 10 0.0050000    11   0.45641 0.53602  0.041593
```

Decision Trees  
ooooo

Regression Trees  
oooooooooooooooooooo

Pruning  
ooooooo

Trees in R  
oooooooo●ooo

## Analyze Results cont'd

- Detailed listing of model parts can be accessed via `summary`:

## Analyze Results cont'd

- Detailed listing of model parts can be accessed via `summary`:

```
summary(tree_model2)
```

```
## Call:  
## rpart(formula = Carbon_Sequestration_lb ~ Tree_Height + Crown_Width,  
##        data = my_pdxTrees, control = rpart.control(minsplit = 20,  
##        xval = 10, maxdepth = 10, cp = 0.005))  
## n=1037 (2 observations deleted due to missingness)  
##  
##          CP nsplit rel error     xerror      xstd  
## 1  0.304594426    0 1.0000000 1.0015819 0.06083997  
## 2  0.127260632    1 0.6954056 0.7296616 0.04066145  
## 3  0.025587347    2 0.5681449 0.6138089 0.03803769  
## 4  0.015177861    3 0.5425576 0.5753989 0.03914593  
## 5  0.014222123    5 0.5122019 0.5705610 0.03960569  
## 6  0.010849075    6 0.4979797 0.5548873 0.03897808  
## 7  0.009024312    7 0.4871307 0.5342982 0.03851474  
## 8  0.006608748    9 0.4690820 0.5269734 0.04016718  
## 9  0.006064592   10 0.4624733 0.5365762 0.04102266  
## 10 0.005000000    11 0.4564087 0.5360229 0.04159291  
##  
## Variable importance  
## Crown_Width Tree_Height  
##          80         20  
##  
## Node number 1: 1037 observations,      complexity param=0.3045944  
## mean=42.47387, MSE=1133.717  
## left son=2 (778 obs) right son=3 (259 obs)  
## Primary splits:  
##   Crown_Width < 56.25 to the left,  improve=0.3025602, (3 missing)  
##   Tree_Height < 42.5 to the left,  improve=0.2366680, (0 missing)  
## Surrogate splits:  
##   Tree_Height < 149.5 to the left,  agree=0.75, adj=0.004, (3 split)
```

Decision Trees  
ooooo

Regression Trees  
oooooooooooooooooooo

Pruning  
ooooooo

Trees in R  
oooooooo●ooo


## CV Plots

- We can plot the results of cross-validation using `plotcp`:

## CV Plots

- We can plot the results of cross-validation using `plotcp`:

```
plotcp(tree_model2)
```



## Pruning

- Based on the CV plot, 10 leaves with  $CP = 0.0077$  gives the lowest error
  - While 7 leaves with  $CP = 0.012$  gives smallest tree within 1 SE of best.



## Pruning

- Based on the CV plot, 10 leaves with  $CP = 0.0077$  gives the lowest error
  - While 7 leaves with  $CP = 0.012$  gives smallest tree within 1 SE of best.
- We can prune our tree using the `prune` function with a given value of `cp`

# Pruning

- Based on the CV plot, 10 leaves with  $CP = 0.0077$  gives the lowest error
  - While 7 leaves with  $CP = 0.012$  gives smallest tree within 1 SE of best.
- We can prune our tree using the `prune` function with a given value of `cp`

```
pruned_tree <- prune(tree_model2, cp = 0.0077)
```



## Test Error Rates

- How well do models do on the test data?

## Test Error Rates

- How well do models do on the test data?

- Let's build a results data frame:

```
results <- data.frame(model = "full",
                      obs = my_pdxTrees_test$Carbon_Sequestration_lb,
                      preds = predict(tree_model2, my_pdxTrees_test))
results <- rbind(results,
                 data.frame(model = "pruned",
                            obs = my_pdxTrees_test$Carbon_Sequestration_lb,
                            preds = predict(pruned_tree, my_pdxTrees_test)))
```

## Test Error Rates

- How well do models do on the test data?

- Let's build a results data frame:

```
results <- data.frame(model = "full",
                      obs = my_pdxTrees_test$Carbon_Sequestration_lb,
                      preds = predict(tree_model2, my_pdxTrees_test))
results <- rbind(results,
                 data.frame(model = "pruned",
                            obs = my_pdxTrees_test$Carbon_Sequestration_lb,
                            preds = predict(pruned_tree, my_pdxTrees_test)))
```

- And use rmse from yardstick to assess:

## Test Error Rates

- How well do models do on the test data?

- Let's build a results data frame:

```
results <- data.frame(model = "full",
                       obs = my_pdxTrees_test$Carbon_Sequestration_lb,
                       preds = predict(tree_model2, my_pdxTrees_test))
results <- rbind(results,
                 data.frame(model = "pruned",
                             obs = my_pdxTrees_test$Carbon_Sequestration_lb,
                             preds = predict(pruned_tree, my_pdxTrees_test)))
```

- And use rmse from yardstick to assess:

```
library(yardstick)
results %>% group_by(model) %>%
  rmse(truth = obs, estimate = preds) %>% arrange(.estimate)

## # A tibble: 2 x 4
##   model .metric .estimator .estimate
##   <chr>  <chr>    <chr>        <dbl>
## 1 pruned rmse    standard     14.2
## 2 full    rmse    standard     15.4
```